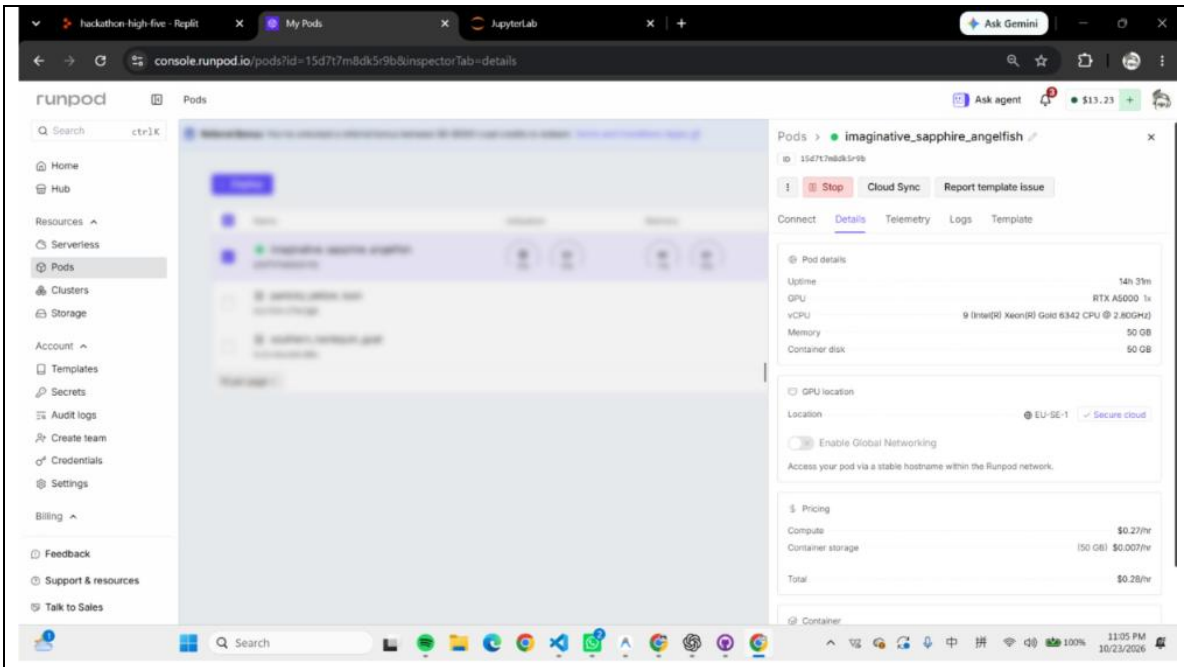
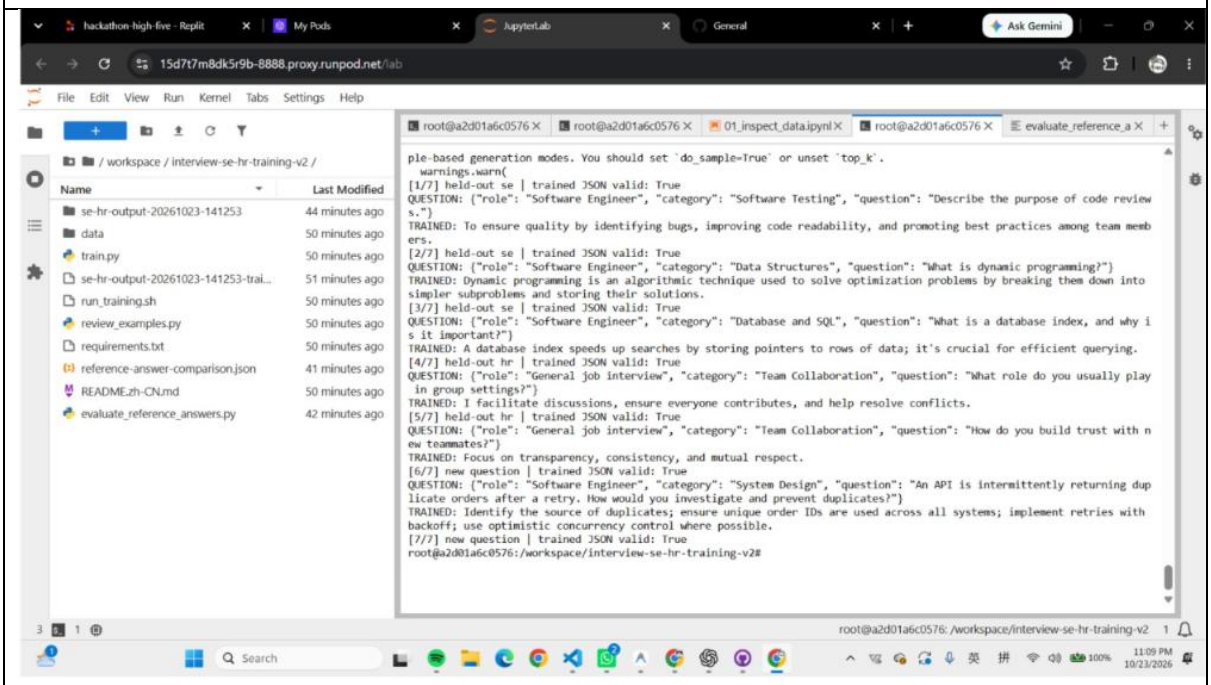
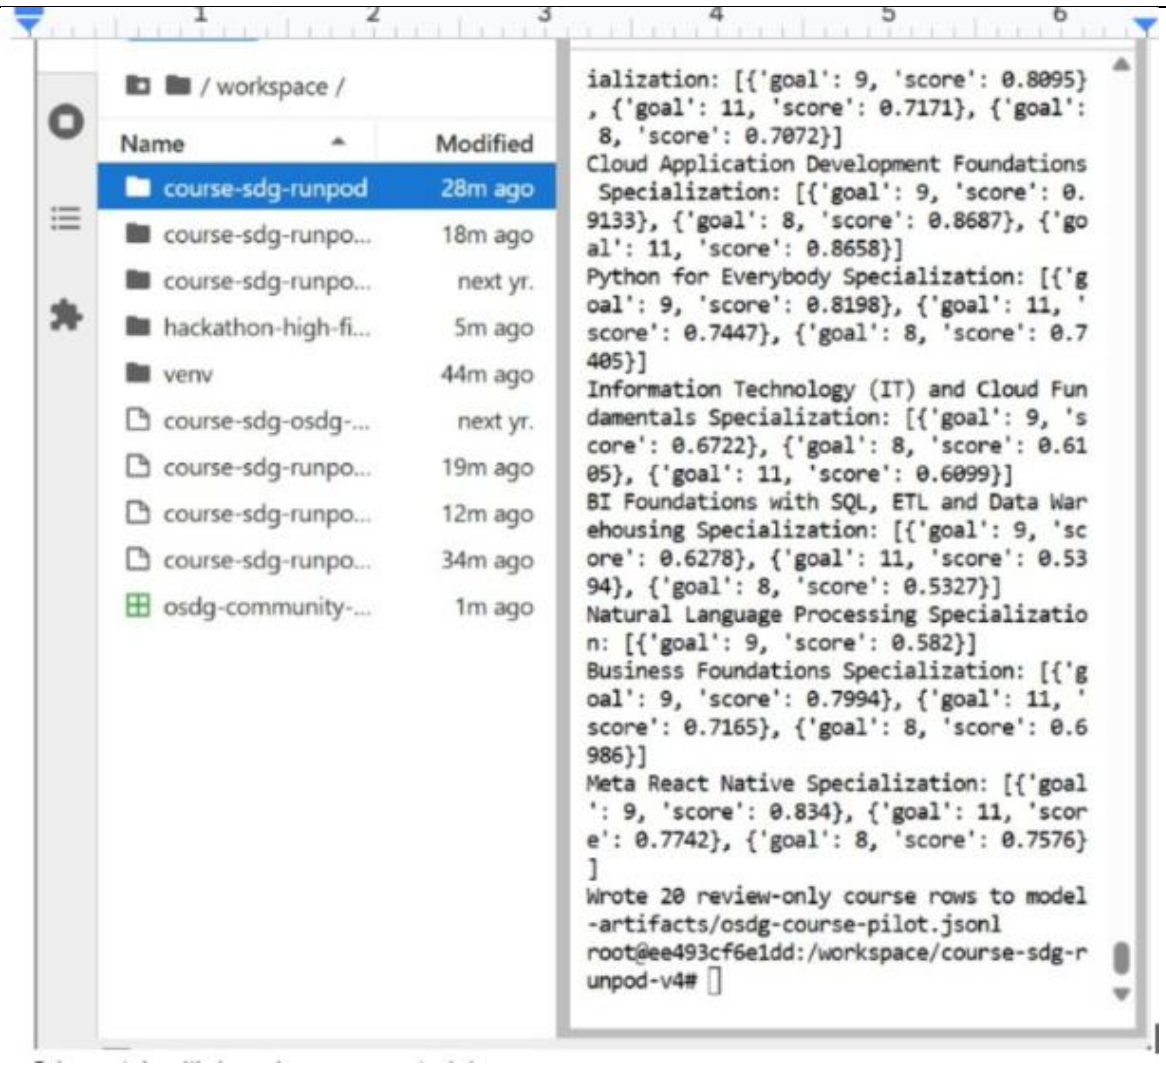
Runpod Prove of Use



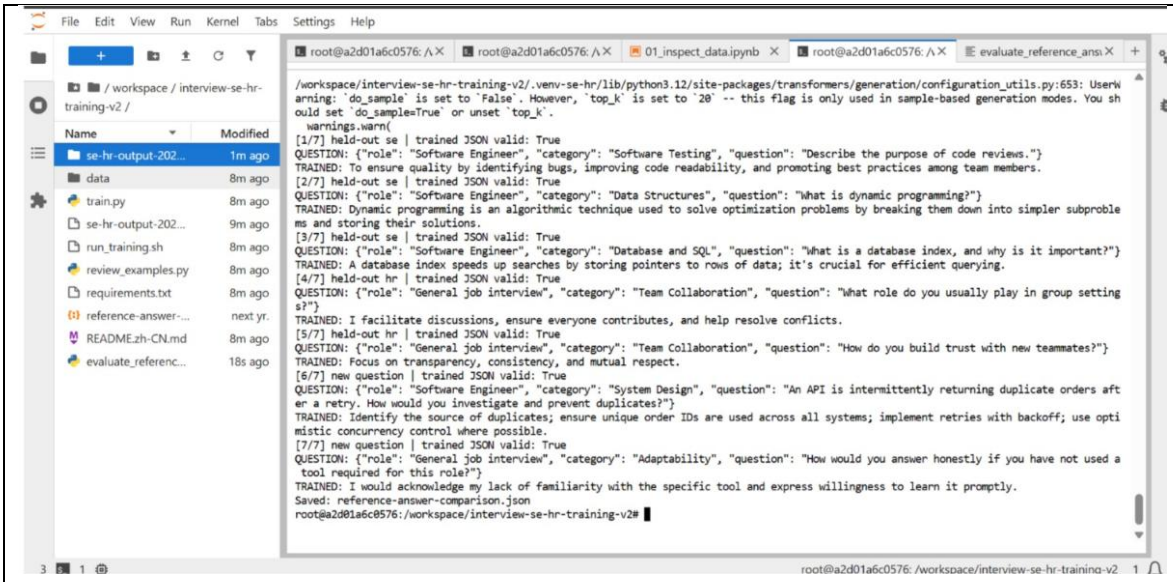
The screenshot shows how we use **RunPod** to manage cloud GPU resources for AI model development and deployment. The platform provides access to an NVIDIA RTX A5000 GPU, allowing us to run computationally intensive AI workloads. It also displays GPU specifications, resource allocation, uptime, and hourly pricing, helping our team monitor performance and manage computational costs efficiently.



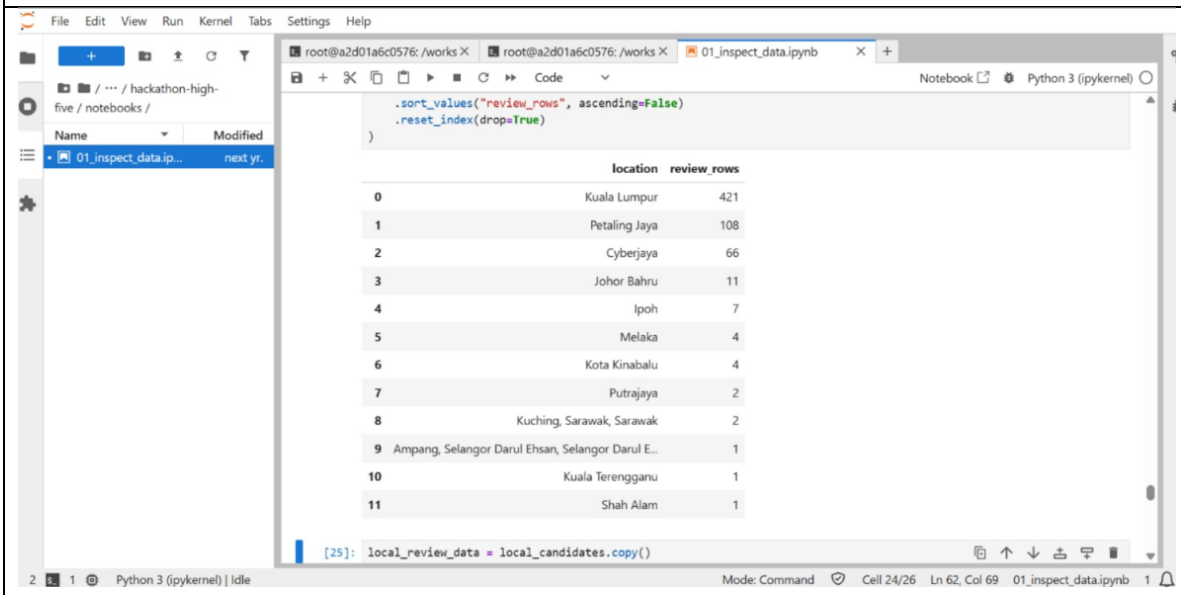
The screenshot shows how we use **JupyterLab on RunPod** to train and evaluate our AI-powered interview question-answering model. We execute Python scripts, process interview datasets, and compare model-generated answers with reference answers. The terminal also displays validation results for 7 interview questions, helping us verify the output format and assess model performance.



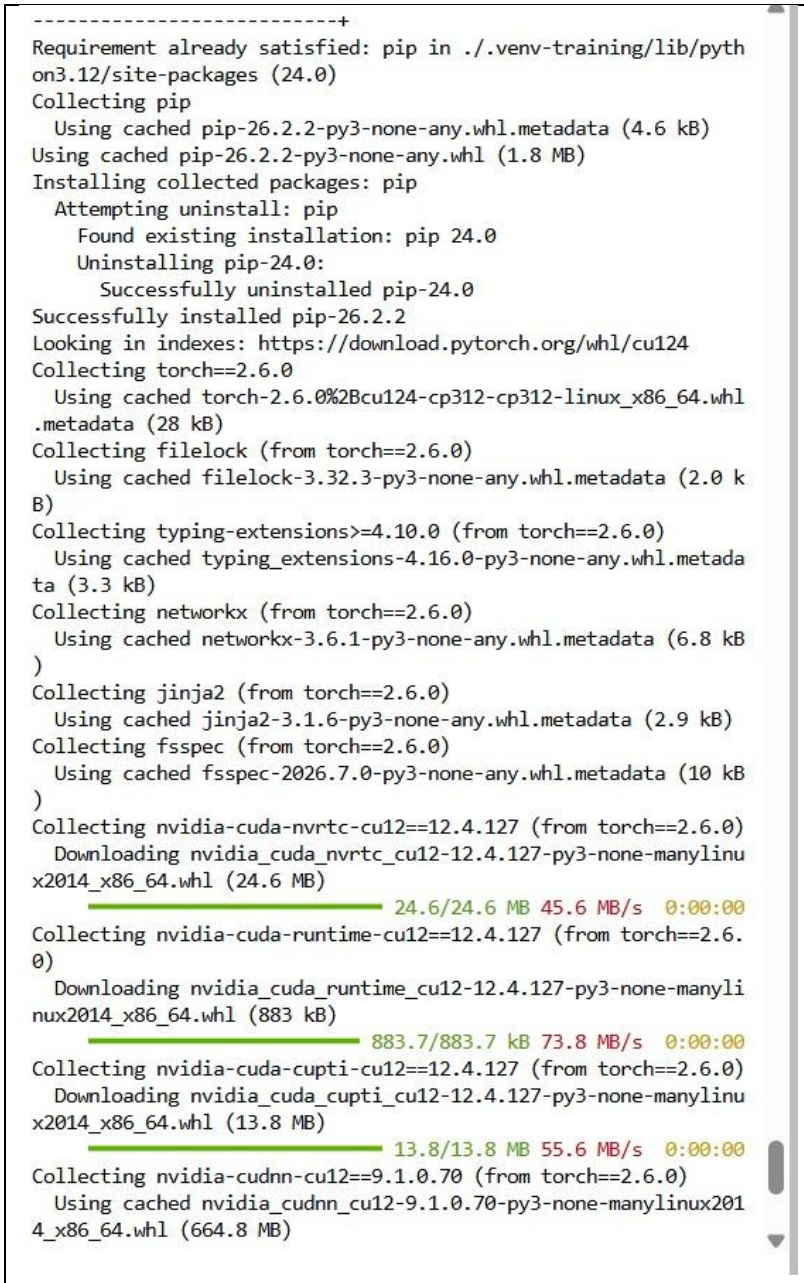
In **RunPod JupyterLab,** it is used to process and evaluate course recommendations based on their alignment with UN Sustainable Development Goals (SDGs). The system generates relevance scores for different courses against SDG goals and saves 20 course evaluation results in a JSONL file for further analysis and validation.



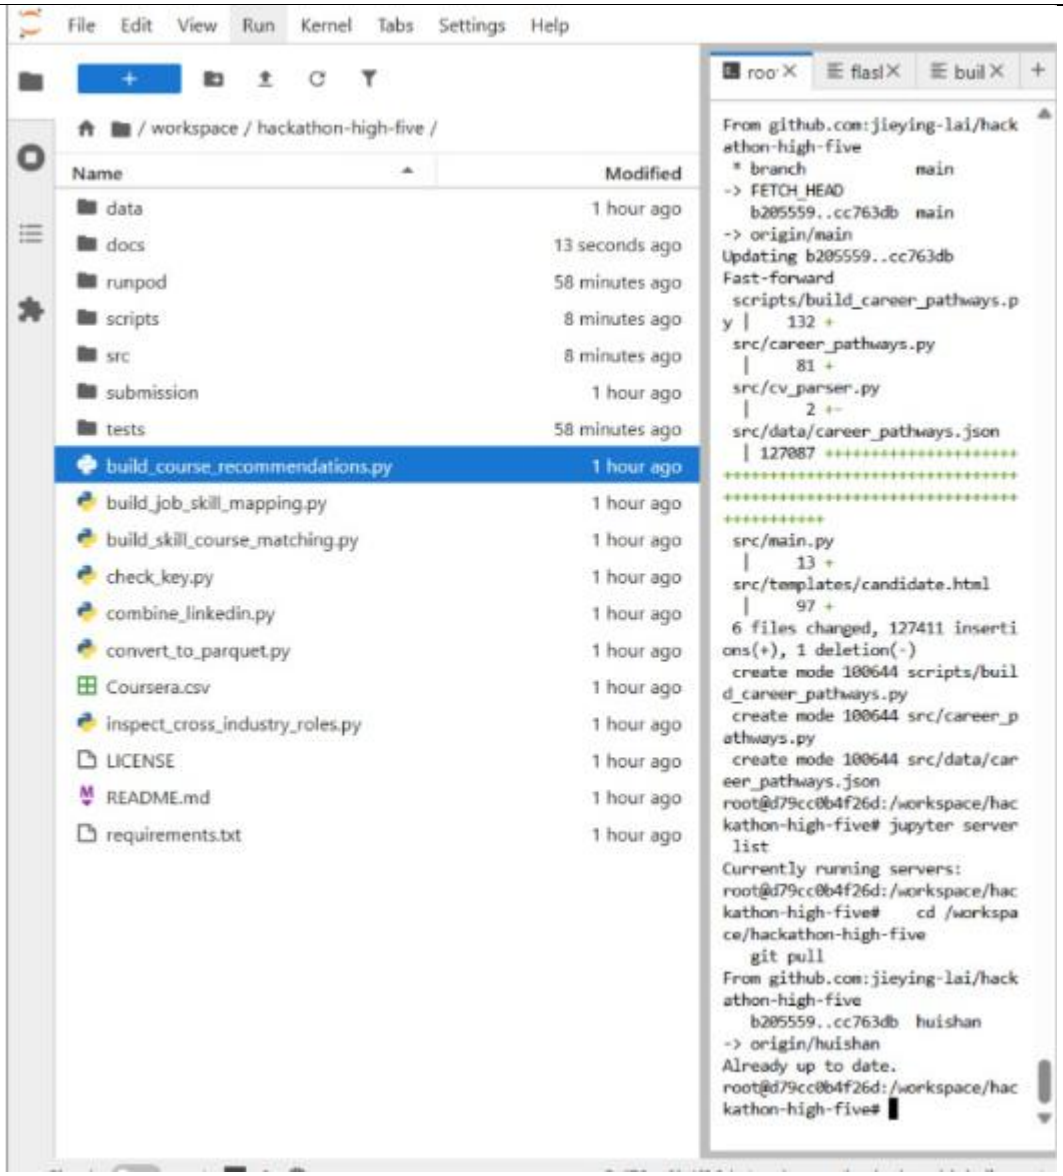
In **RunPod JupyterLab,** it is utilised to test and validate our AI-powered interview answer generation model. The system generates answers for technical and behavioural interview questions, checks the JSON output format, and compares generated responses with reference answers. All seven test cases passed JSON format validation, and the comparison results were saved for further evaluation.

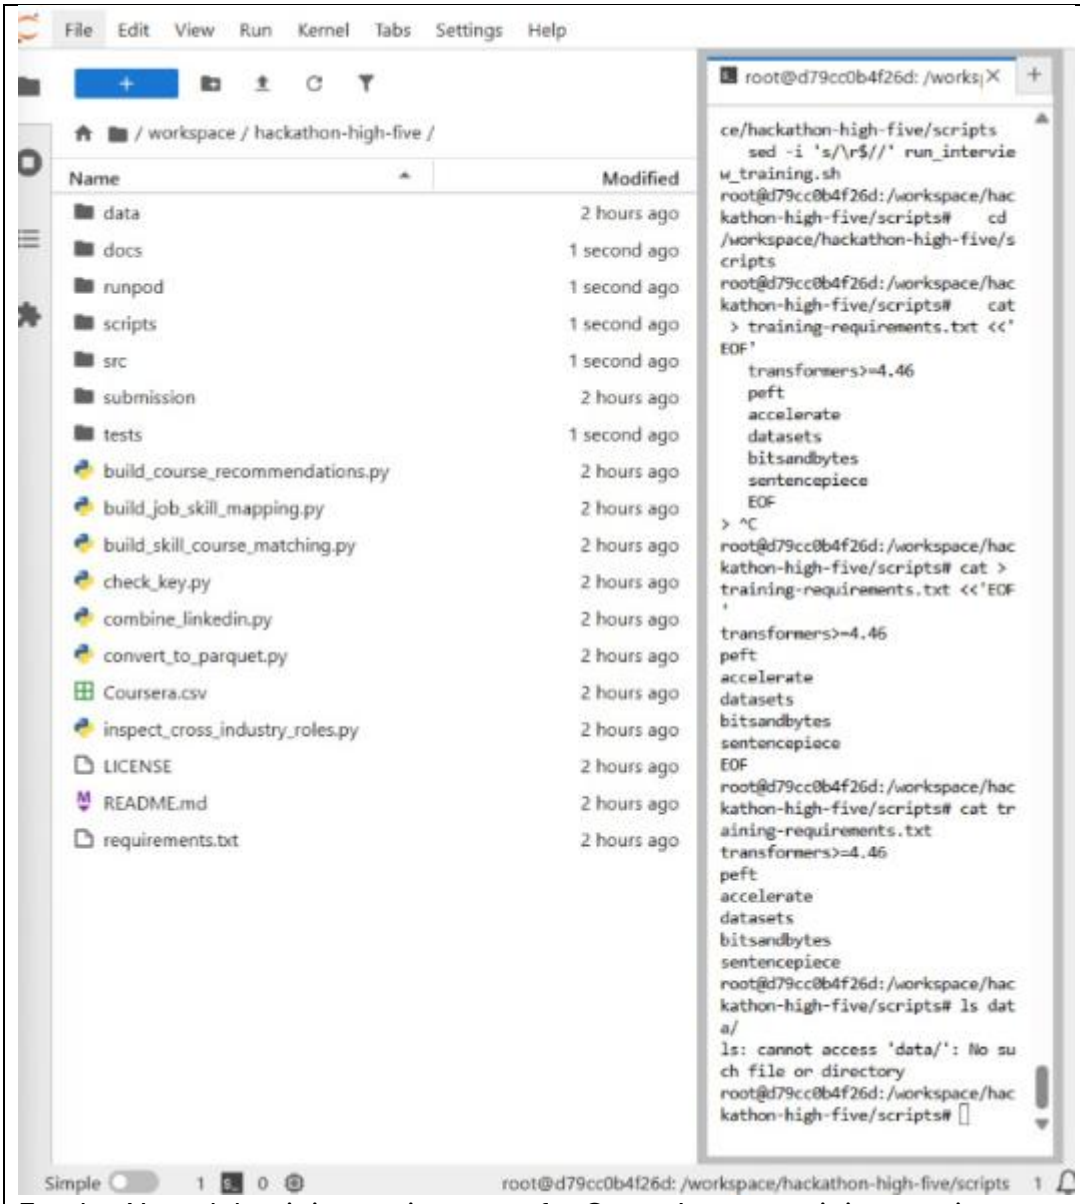


In **JupyterLab on RunPod,** company review data is applied to explore and analyse using Python. We group company reviews by location, calculate the number of reviews in each area, and identify regions with higher data availability. This helps us understand the dataset's geographical distribution and assess its suitability for our company's review analysis feature.

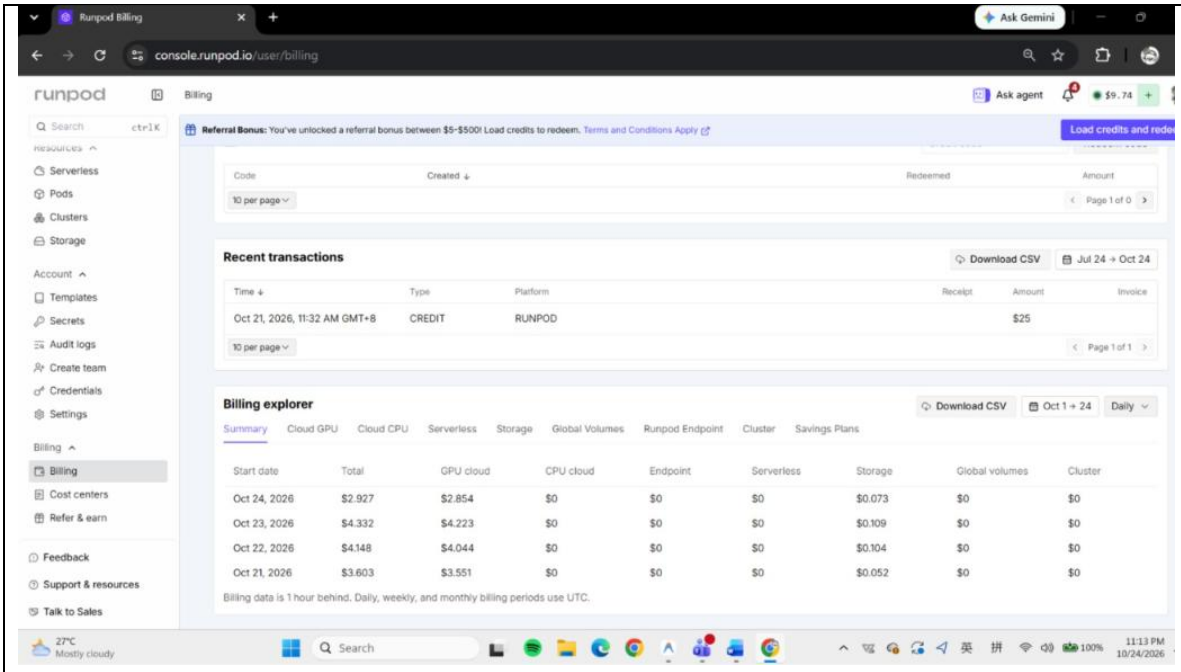


Set up the AI model training environment on **RunPod** by upgrading pip and installing PyTorch with CUDA 12.4 support. This enables GPU-accelerated computation for AI model training and inference, helping improve processing efficiency when handling computationally intensive tasks.

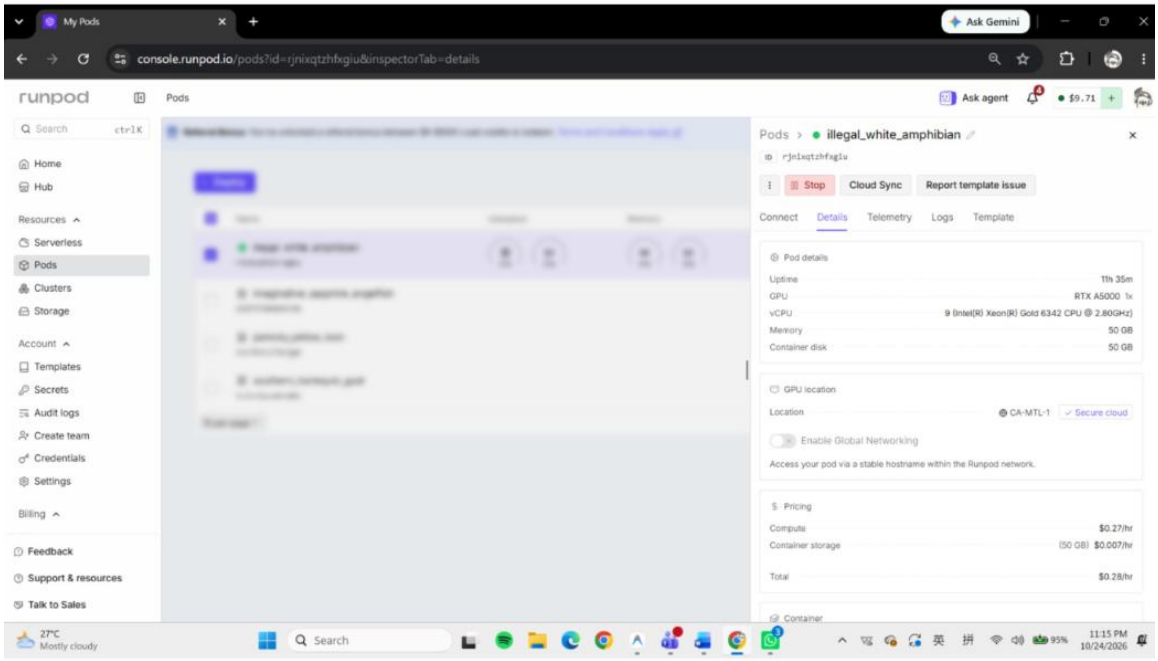

File Edit View Run Kernel Tabs Settings Help
/ workspace / hackathon-high-five /
Name
Modified
data 1 hour ago
docs 13 seconds ago
runpod 58 minutes ago
scripts 8 minutes ago
src 8 minutes ago
submission 1 hour ago
tests 58 minutes ago
build_course_recommendations.py 1 hour ago
build_job_skill_mapping.py 1 hour ago
build_skill_course_matching.py 1 hour ago
check_key.py 1 hour ago
combine_linkedin.py 1 hour ago
convert_to_parquet.py 1 hour ago
Coursera.csv 1 hour ago
inspect_cross_industry_roles.py 1 hour ago
LICENSE 1 hour ago
README.md 1 hour ago
requirements.txt 1 hour ago
roo
flas
buil
From github.com:jieying-lai/hackathon-high-five
* branch main
-> FETCH_HEAD
b205559..cc763db main
-> origin/main
Updating b205559..cc763db
Fast-forward
scripts/build_career_pathways.py | 132 +
src/career_pathways.py | 81 +
src/cv_parser.py | 2 +-
src/data/career_pathways.json | 127087 ++++++



For the AI model training environment for CareerLens, a training requirements file is created with containing essential Python libraries, including Transformers, PEFT, Accelerate, Datasets, BitsAndBytes, and SentencePiece. These dependencies support model fine-tuning, efficient GPU utilisation, and dataset processing.



Billing page



GPU usage